
Group Homework3 4 :

Michelle (Chia Chi) Ho, Jiejun Lu, Jiawen Tong

Homework 3

Data: Homework_3_Data.txt, housedata.zip

Harvard University

Fall 2018

Instructors: Rahul Dave

Due Date: Saturday, September 29th, 2018 at 11:59pm

Instructions:

- Upload your final answers in the form of a Jupyter notebook containing all work to Canvas.
- Structure your notebook and your work to maximize readability.

```
In [1]: 1 import numpy as np
2 import scipy.stats
3 import scipy.special
4
5 import matplotlib
6 import matplotlib.pyplot as plt
7 import matplotlib.mlab as mlab
8 from matplotlib import cm
9 import pandas as pd
10 import seaborn as sns
11 %matplotlib inline
12
13 from scipy.stats import norm
```

Question 1: When have no confidence that you can lift yourself by the Bootstrap?

Coding required

The idea behind non-parametric bootstrapping is that sampling distributions constructed via the true data generating process should be very close to sampling distributions constructed by resampling. We mentioned in lab that one edge cases for bootstrapping is calculating order statistics. Let's explore this edgecase.

- 1.1. Suppose you have $\{X_1, X_2, \dots, X_n\}$ datapoints such that X_i are independently and identically drawn from a $Unif(0, \theta)$. Consider the extreme order statistic $Y = X_{(n)} = \max(X_1, X_2, \dots, X_n)$. Write an expression for the distribution $f_Y(Y|\theta)$.
- 1.2. Derive $\hat{\theta}$ the maximum likelihood estimate for θ given datapoints $\{X_1, X_2, \dots, X_n\}$.
- 1.3. To see an alternate potential estimator use the distribution you derived in 1.1. to find an expression for the unbiased estimate of theta.
- 1.4. Use scipy/numpy to generate 100 samples $\{X_i\}$ from $Unif(0,1)$ (i.e. let $\theta = 1$) and store them in the variable `original_xi_samples`. Based on your data sample, what's the empirical estimate for θ .
- 1.6. Use non-parametric bootstrap to generate a sampling distribution of 1000 estimates for theta. Plot a histogram of your sampling distribution. Make sure to title and label the plot.
- 1.7. Is your histogram smooth? From visual inspection does it seem like a good representation of a sampling distribution?
- 1.8. So far we've used a "natural" version of calculating bootstrap confidence intervals -- the percentile method. In this situation is it possible for the "true" value of θ to be in the confidence interval? In order to remedy this we'll use an alternate confidence interval version called the pivot confidence interval. The pivot confidence interval is defined as $[2\hat{\theta} - \hat{\theta}_{(0.975)}^*, 2\hat{\theta} - \hat{\theta}_{(0.025)}^*]$. Is the true value contained in this interval?

1.1 Answer

The CDF of Y is:

$$F(y|\theta) = \int_{-\infty}^y f_Y(y|\theta) dy = P(Y \leq y) = P(\max(X_{1:n}) \leq y) = \begin{cases} 0 & y < 0 \\ (\frac{y}{\theta})^n & 0 \leq y \leq \theta \\ 1 & y > \theta \end{cases}$$

Therefore, the PDF of Y would be:

$$\Rightarrow f_Y(y|\theta) = \frac{\partial F_Y(y|\theta)}{\partial y} = \begin{cases} 0 & y < 0 \\ \frac{ny^{n-1}}{\theta^n} & 0 \leq y \leq \theta \\ 0 & y > \theta \end{cases}$$

1.2 Answer

Let $\max(X_{1:n}) = m$; we need to maximize the likelihood subject to $0 \leq m \leq \theta$:

$$\text{Likelihood} = \prod_i^n P(X_i = x_i) = \prod_i^n P(X_i < \max(X_{1:n})) = \left(\frac{m}{\theta}\right)^n$$

$$\log(\text{Likelihood}) = n \log m - n \log \theta$$

$$\frac{\partial \log(\text{Likelihood})}{\partial \theta} = -\frac{n}{\theta} < 0$$

Therefore, the log likelihood is decreasing for any $\theta > 0$ and the maximum likelihood is at θ 's minimum such that $\theta \geq \max(X_1, X_2, \dots, X_n)$:

$$\hat{\theta}_{MLE} = \max(X_1, X_2, \dots, X_n) = Y$$

1.3 Answer

$$E[Y] = \int y f_Y(y|\theta) dy = \int_0^\theta y \frac{ny^{n-1}}{\theta^n} dy = \frac{n}{\theta^n} \int_0^\theta y^n dy = \frac{n}{n+1} \theta$$

$$\Rightarrow \theta = \frac{n+1}{n} E[Y]$$

Since

$$E[\theta] = \frac{n+1}{n} E[Y] = \theta$$

, this is an unbiased estimate for θ .

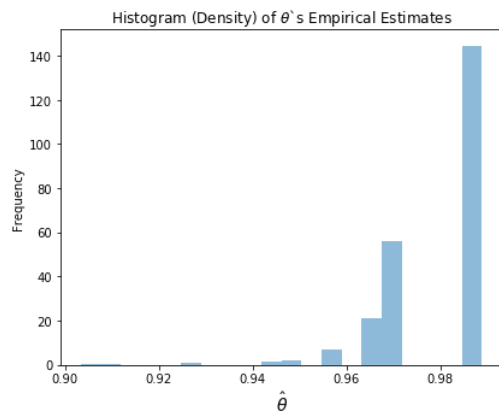
1.4 Answer

```
In [2]: 1 # 1.4
        2 np.random.seed(1)
        3
        4 theta = 1.
        5 n_size = 100
        6 original_xi_samples = np.random.uniform(0, theta, size=n_size)
        7 theta_hat = max(original_xi_samples)
        8 print('Empirical Estimate of theta = max(original_xi_samples):', theta_hat)
```

Empirical Estimate of theta = max(original_xi_samples): 0.9888610889064947

1.6 Answer

```
In [3]: 1 # 1.6
        2 np.random.seed(1)
        3
        4 n_sim = 1000
        5 theta_estimates = []
        6 for i in range(n_sim):
        7     # extract the empirical theta on each bootstrapped sample
        8     theta_estimates.append(max(np.random.choice(original_xi_samples, size=n_size, replace=True)))
        9
        10 ax = pd.Series(theta_estimates, name='Empirical Estimate of theta').plot(kind='hist', density=True, bins=20,
        11     figsize=(6, 5), alpha=0.5, title=r'Histogram (Density) of $\hat{\theta}$'s Empirical Estimates')
        12 ax.set_xlabel(r'$\hat{\theta}$', fontsize=14)
        13 plt.tight_layout()
```



1.7 Answer

The histogram is not smooth and thus does not seem like a good representation of the sampling distribution of θ . The discontinuity is due to the non-parametric bootstrapping.

```
In [4]: 1 # 1.8
2 r = np.percentile(theta_estimates, 97.5)
3 l = np.percentile(theta_estimates, 2.5)
4 print('--- Natural 95% CI ---')
5 print('left: {}, right: {}'.format(l, r))
6 print('True theta is in the natural CI:', (l<=theta) and (theta<=r))
7
8 l_new = 2*theta_hat - r
9 r_new = 2*theta_hat - l
10 print('\n--- Modified 95% CI ---')
11 print('left: {}, right: {}'.format(l_new, r_new))
12 print('True theta is in the modified CI:', (l_new<=theta) and (theta<=r_new))
```

```
--- Natural 95% CI ---
left: 0.9578895301505019, right: 0.9888610889064947
True theta is in the natural CI: False
```

```
--- Modified 95% CI ---
left: 0.9888610889064947, right: 1.0198326476624875
True theta is in the modified CI: True
```

1.8 Answer

As shown above, the true θ is not contained in the natural CI but is contained in the modified one.

Question 2: Visualize Your Poor Marginalized Conditional Love

Coding required

Read the data set contained in [data/Homework_3_Data.txt](#) ([data/Homework_3_Data.txt](#)). Each data point is a two-dimensional vector, $\mathbf{x} = (x_1, x_2)$.

- 2.1. Make a 2-D visualization of the empirical distribution of the data.
- 2.2. We assume that the data was generated by some probability distribution (pdf). Visualize that pdf, f_X .
- 2.3. Visualize the conditional distribution defined by $f_{x_2|x_1}$ for $x_1 \in [3.99, 4.01]$.
- 2.4. Visualize the marginal distribution defined by f_{x_1} .

2.1~2.4 Answer

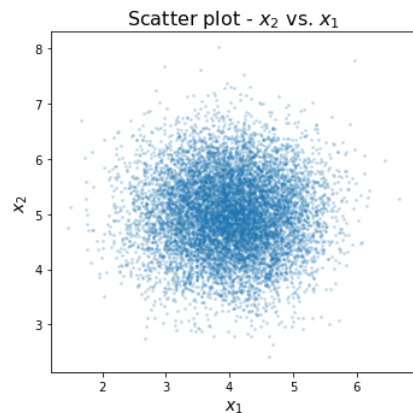
```
In [5]: 1 # read in the data
2 data = pd.read_table('data/Homework_3_Data.txt', sep=',', header=None)
3 data.rename(columns={0:'x1', 1:'x2'}, inplace=True)
4 print(data.shape)
5 data.head()
```

```
(10000, 2)
```

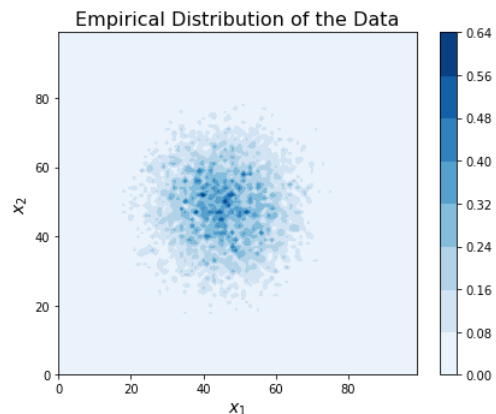
```
Out[5]:
```

	x1	x2
0	3.008992	6.205285
1	4.845897	4.864804
2	5.137567	4.536671
3	4.766038	4.158884
4	4.242169	4.070555

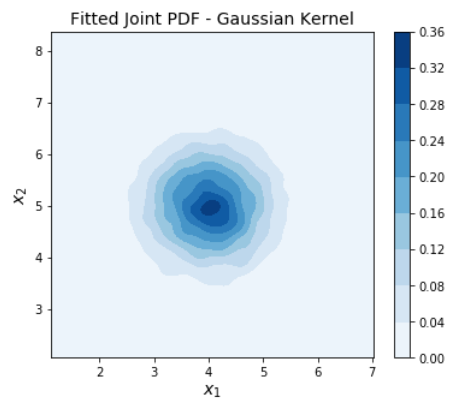
```
In [6]: 1 # 2D scatter plot of the data
2 fig, ax = plt.subplots(1, 1, figsize=(5, 5))
3 ax.scatter(x=data['x1'], y=data['x2'], alpha=0.2, s=3)
4
5 ax.set_title(r'Scatter plot -  $x_2$  vs.  $x_1$ ', fontsize=16)
6 ax.set_ylabel(r' $x_2$ ', fontsize=14)
7 ax.set_xlabel(r' $x_1$ ', fontsize=14)
8 plt.tight_layout()
```



```
In [7]: 1 # 2.1 Visualize the empirical distribution of the 2D data in Contourf Plot
2 H, _, _ = np.histogram2d(x=data['x1'], y=data['x2'], bins=100, normed=True)
3 plt.subplots(figsize=(6, 5))
4 plt.contourf(H, cmap='Blues')
5 plt.colorbar()
6 plt.title('Empirical Distribution of the Data', fontsize=16)
7 plt.xlabel(r' $x_1$ ', fontsize=14)
8 plt.ylabel(r' $x_2$ ', fontsize=14)
9 plt.tight_layout()
```



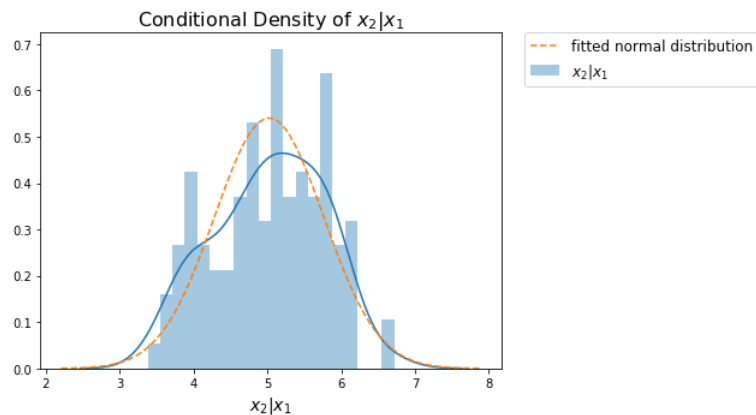
```
In [8]: 1 # 2.2 Visualize the fitted 2D joint pdf
2 plt.subplots(1, 1, figsize=(6, 5))
3 sns.kdeplot(data['x1'], data['x2'], shade=True, kernel='gau', cmap='Blues', cbar=True)
4 plt.title('Fitted Joint PDF - Gaussian Kernel', fontsize=14)
5 plt.xlabel(r' $x_1$ ', fontsize=14)
6 plt.ylabel(r' $x_2$ ', fontsize=14)
7 plt.show()
```



```

In [9]: 1 # 2.3 Visualize the conditional distribution defined by  $f_{x_2|x_1}$  for  $x_1$  in [3.99,4.01]
2 plt.subplots(1, 1, figsize=(6, 5))
3 sns.distplot(data[(data['x1']>=3.99) & (data['x1']<=4.01)]['x2'], bins=20, hist=True, kde=True, label=r'$x_2|x_1$')
4
5 # fit a normal distribution to 'Ash'
6 mu, std = norm.fit(data[(data['x1']>=3.99) & (data['x1']<=4.01)]['x2'].values)
7 # plot the fitted normal distribution
8 xmin, xmax = plt.xlim()
9 x = np.linspace(xmin, xmax, 100)
10 p_drop = norm.pdf(x, mu, std)
11 plt.plot(x, p_drop, '--', label='fitted normal distribution')
12
13 plt.title(r'Conditional Density of  $x_2|x_1$ ', fontsize=16)
14 plt.xlabel(r'$x_2|x_1$', fontsize=14)
15 plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0., fontsize=12)
16 plt.tight_layout()

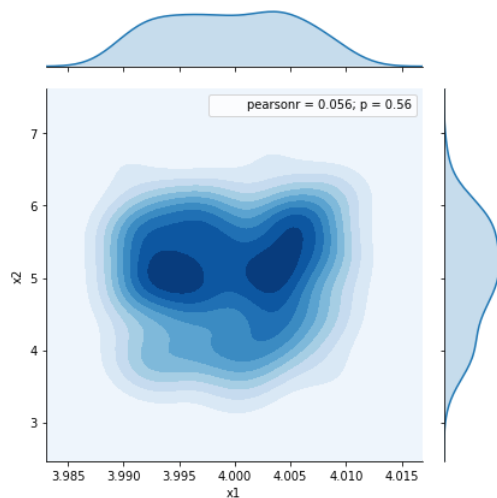
```



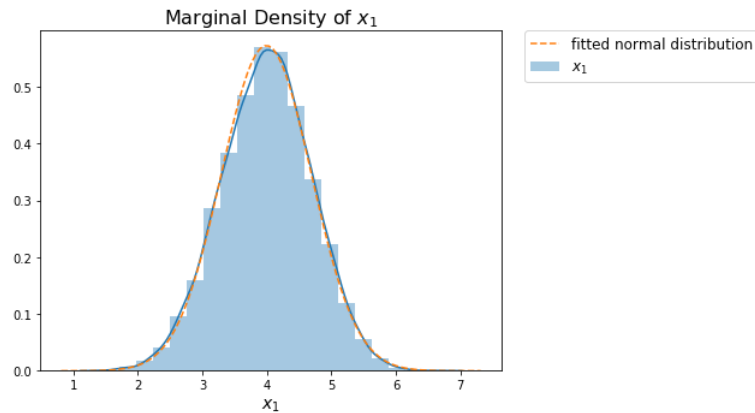
```

In [10]: 1 # 2.3 Visualize the conditional distribution defined by  $f_{x_2|x_1}$  for  $x_1$  in [3.99,4.01]
2 # if we do not take  $x_1$  as constant
3 sns.jointplot('x1', 'x2', data=data[(data['x1']>=3.99) & (data['x1']<=4.01)], kind='kde', cmap='Blues')
4 plt.tight_layout()

```



```
In [11]: 1 # 2.4 Visualize x1's marginal distribution f_{x1}
2 plt.subplots(1, 1, figsize=(6, 5))
3 sns.distplot(data['x1'], bins=20, label=r'$x_1$')
4
5 # fit a normal distribution
6 mu, std = norm.fit(data['x1'].values)
7 # plot the fitted normal distribution
8 xmin, xmax = plt.xlim()
9 x = np.linspace(xmin, xmax, 100)
10 p_drop = norm.pdf(x, mu, std)
11 plt.plot(x, p_drop, '--', label='fitted normal distribution')
12
13 plt.title(r'Marginal Density of $x_1$', fontsize=16)
14 plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0., fontsize=12)
15 plt.xlabel(r'$x_1$', fontsize=14)
16 plt.tight_layout()
```



2.5. Empirically estimate the mean of the distribution f_{x_1} . Estimate, also the SE (standard error) of the estimate.

2.6. Empirically estimate the standard deviation of the distribution $f_{x_2|x_1}$, for $x_1 \in [3.99, 4.01]$. Estimate, also the SE (standard error) of the estimate.

2.7. Given the SE, How many digits in your standard deviation estimate are significant? Explain why.

In obtaining estimates for this problem we want you to

- define a function called `get_bootstrap_sample(dataset)` to generate each bootstrap sample
- and then another function `perform_bootstrap(dataset)` to generate all the samples.

They should both take as parameters the dataset from which you'll be drawing samples. `perform_bootstrap` should call `get_bootstrap_sample` and return a sequence of bootstrap samples. `get_bootstrap_sample` should return an individual bootstrap sample.

2.5~2.7 Answer

```
In [12]: 1 def get_bootstrap_sample(dataset):
2         return np.random.choice(dataset, size=len(dataset), replace=True)
3
4 def perform_bootstrap(dataset, n_sim=1000):
5     all_samples = []
6     for i in range(n_sim):
7         all_samples.append(get_bootstrap_sample(dataset))
8     return all_samples
9
```

```
In [13]: 1 # 2.5 Empirically estimate the mean of f_{x1}. Estimate, also the SE (standard error) of the estimate.
2 np.random.seed(1)
3
4 sample_x1 = perform_bootstrap(data['x1'].values)
5 x1_sample_mean_estimates = np.mean(sample_x1, axis=1)
6 print('The mean of f_{x1}\s mean estimates is:', np.mean(x1_sample_mean_estimates))
7 print('with standard error:', np.std(x1_sample_mean_estimates))
```

The mean of f_{x_1} 's mean estimates is: 3.9927692743164243
with standard error: 0.006865925559635069

```
In [14]: 1 # 2.6 Empirically estimate the std of  $f_{x_2|x_1}$ , for  $x_1$  in [3.99, 4.01]
2 # Estimate, also the SE (standard error) of the estimate.
3 data_subset = data[(data['x1']>=3.99)&(data['x1']<=4.01)].reset_index(drop=True)
4
5 # generate bootstrapped samples where  $x_1$  is in [3.99, 4.01]
6 np.random.seed(1)
7 sample_x2_given_x1 = perform_bootstrap(data[(data['x1']>=3.99)&(data['x1']<=4.01)]['x2'].values)
8 x2_given_x1_std_estimates = np.std(sample_x2_given_x1, axis=1)
9 print('The mean of  $f_{x_2|x_1}$ 's std estimates is:', np.mean(x2_given_x1_std_estimates))
10 print('with standard error:', np.std(x2_given_x1_std_estimates))
```

The mean of $f_{x_2|x_1}$'s std estimates is: 0.7329488865102515
with standard error: 0.03895510137143677

2.7 Answer

The standard deviation estimate of conditional pdf $f_{x_2|x_1} \approx 0.733$ with a SE=0.03+. This suggests that the last digit worth reporting in the estimate is the second decimal place as it is the first digit to encapsulate the error. Therefore, the estimate has 2 significant digits.

Similarly, the mean estimate of marginal pdf $f_{x_1} \approx 3.993$ with a SE=0.006+. This suggests that the last digit worth reporting in the estimate is the third decimal place. Therefore, the estimate has 4 significant digits.

Problem 3: Linear Regression

Consider the following base Regression class, which roughly follows the API in the python package `scikit-learn`.

Our model is the the multivariate linear model whose MLE solution or equivalent cost minimization was talked about in lecture:

$$y = X\beta + \epsilon$$

where y is a length n vector, X is an $m \times p$ matrix created by stacking the features for each data point, and β is a p length vector of coefficients.

The class showcases the API:

fit(X, y): Fits linear model to X and y .

get_params(): Returns $\hat{\beta}$ for the fitted model. The parameters should be stored in a dictionary with keys "intercept" and "coef" that give us $\hat{\beta}_0$ and $\hat{\beta}_1$. (The second value here is thus a numpy array of coefficient values)

predict(X): Predict new values with the fitted model given X .

score(X, y): Returns R^2 value of the fitted model.

set_params(): Manually set the parameters of the linear model.

```
In [15]: 1 class Regression(object):
2
3     def __init__(self):
4         self.params = dict()
5
6     def get_params(self, k):
7         return self.params[k]
8
9     def set_params(self, **kwargs):
10        for k,v in kwargs.items():
11            self.params[k] = v
12
13    def fit(self, X, y):
14        raise NotImplementedError()
15
16    def predict(self, X):
17        raise NotImplementedError()
18
19    def score(self, X, y):
20        raise NotImplementedError()
```

3.1. In a jupyter notebook code cell below we've defined and implemented the class `Regression`. Inherit from this class to create an ordinary least squares Linear Regression class called `AM207OLS`. Your class will implement an sklearn-like api. It's signature will look like this:

```
class OLS(Regression):
```

Implement `fit`, `predict` and `score`. This will involve some linear algebra. (You might want to read up on pseudo-inverses before you directly implement the linear algebra on the lecture slides).

The R^2 score is defined as:

$$R^2 = 1 - \frac{SS_E}{SS_T}$$

Where:

$$SS_T = \sum_i (y_i - \bar{y})^2, SS_R = \sum_i (\hat{y}_i - \bar{y})^2, SS_E = \sum_i (y_i - \hat{y}_i)^2$$

where y_i are the original data values, \hat{y}_i are the predicted values, and \bar{y} is the mean of the original data values.

3.1 Answer - Codes

```
In [16]: 1 class OLS(Regression):
2         def __init__(self):
3             self.params = dict()
4
5         def get_params(self, k):
6             if k not in self.params.keys():
7                 raise Exception('The OLS model must be fitted before calling get_params.')
8             return super(OLS, self).get_params(k)
9
10        def fit(self, X, y):
11            # add constant terms
12            ones_col = np.ones((X.shape[0], 1))
13            X_ones = np.concatenate((ones_col, X), axis=1)
14
15            # calculate betas
16            betas = np.dot(np.linalg.inv(np.dot(X_ones.T, X_ones)), np.dot(X_ones.T, y))
17
18            # save the fitted intercept and coefficients to the params dictionary
19            self.params['intercept'] = betas[0]
20            self.params['coef'] = betas[1:]
21
22            # save the training R2 score
23            self.rsquared = self.score(X, y)
24
25
26        def predict(self, X):
27            # add constant terms
28            ones_col = np.ones((X.shape[0], 1))
29            X_ones = np.concatenate((ones_col, X), axis=1)
30
31            # extract betas from self.params dictionary
32            betas = np.zeros((X_ones.shape[1],))
33            betas[0] = self.params['intercept']
34            betas[1:] = self.params['coef']
35
36            # calculate and return predictions
37            return np.dot(X_ones, betas)
38
39        def score(self, X, y):
40            # predict with X
41            yhat = self.predict(X)
42            # calculate SSE
43            SSE = np.sum(np.square(y - yhat))
44            # calculate SST
45            ybar = np.mean(y)
46            SST = np.sum(np.square(y - ybar))
47            # calculate and return rsquared
48            return 1 - (SSE/SST)
```

3.2. We'll create a synthetic data set using the code below. (Read the documentation for `make_regression` to see what is going on).

Verify that your code recovers these coefficients approximately on doing the fit. Plot the predicted \hat{y} against the actual y . Also calculate the score using the same sets x and y . The usage will look something like:

```
lr = OLS()
lr.fit(X,y)
lr.get_params['coef']
lr.predict(X,y)
lr.score(X,y)
```

```
In [17]: 1 from sklearn.datasets import make_regression
2         import numpy as np
3         np.random.seed(99)
4         X, y, coef = make_regression(n_samples=30, n_features=10, n_informative=10, bias=1, noise=2, coef=True)
5         coef
```

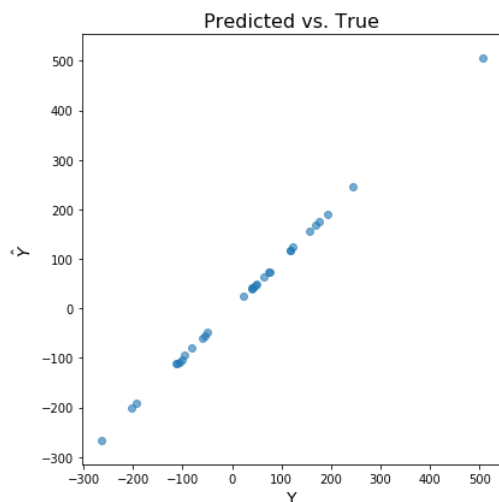
```
Out[17]: array([76.6568183 , 77.67682678, 63.78807738, 19.3299907 , 59.01638708,
53.13633737, 28.77629958, 10.01888939,  9.25346811, 59.55220395])
```

3.2 Answer - Codes


```
In [18]: 1 # 3.2
2 lr = OLS()
3 lr.fit(X,y)
4 print("fitted intercept: ", lr.get_params('intercept'))
5 print("fitted coeffs: ", lr.get_params('coef'))
6 print("=====")
7 print("predictions: ", lr.predict(X))
8 print("=====")
9 print("score: ", lr.score(X,y))

fitted intercept: 1.3458035707838505
fitted coeffs: [77.20719705 76.51004831 62.97865316 18.4436452 58.50019885 53.25126559
28.29088241 9.33333359 10.29584457 59.1606719 ]
=====
predictions: [ 48.57564537 24.85508406 246.39920911 64.72282184 124.00042911
-266.57653702 118.15510334 -108.57077603 191.15229644 174.74404249
-103.59066227 -59.1576374 -54.70947468 73.91003582 505.22781806
39.53820436 -191.02175593 -201.71787963 46.1500923 -111.90307749
117.38777883 -111.52335297 41.66543625 168.33074858 73.41934029
-80.64319083 155.12182695 -94.18157131 -48.17239883 40.78548557]
=====
score: 0.9999155832062194
```

```
In [19]: 1 # plot predicted y against actual y
2 plt.figure(figsize=(6, 6))
3 plt.scatter(x=y, y=lr.predict(X), alpha=0.6)
4
5 plt.title('Predicted vs. True', fontsize=16)
6 plt.ylabel(r'$\hat{Y}$', fontsize=14)
7 plt.xlabel(r'$Y$', fontsize=14)
8 plt.axis('equal')
9 plt.tight_layout()
```



Question 4: Is the Incumbent of the House in?

We shall consider US House data from 1896 to 1990. This dataset was compiled for [Gelman, Andrew, and Gary King, "Estimating incumbency advantage without bias," American Journal of Political Science \(1990\): 1142-1164](http://gking.harvard.edu/files/gking/files/inc.pdf). Why incumbency and why the house? The house gives us lots of races in any given year to validate our model, and in elections which happen every two years, where demography hasn't changed much, incumbency is a large effect, as might be the presence of a national swing (which we would capture in an intercept in a regression).

Let us, then, imagine a very simplified model in which the democratic party's fraction of the vote in this election, for seat(county) i , at time t years, $d_{i,t}$, is a linear combination of the democratic party's fraction of the vote in the previous election, at time $t-2$, $d_{i,t-2}$, and a categorical variable $I_{i,t}$, which characterizes the nature of the candidate running in this election:

$$I = \begin{cases} -1 & \text{Republican Incumbent Running} \\ 0 & \text{New Candidate Running} \\ 1 & \text{Democratic Incumbent Running} \end{cases}$$

We use the statsmodels formula notation:

DP1 ~ DP + I .

This means linear regress DP, the democratic fraction of the vote this time around for a given house seat on DP1 which is the democratic fraction the previous time around and I, a "factor" or categorical(nominal) variable with 3 levels.

In mathematical notation this regression is:

$$d_{i,t} = \beta_1 d_{i,t-2} + \beta_2 I_{i,t} + \beta_0,$$

where $d_{i,t-2}$ is the democratic fraction in county i at the previous election, and $I_{i,t}$ is the factor above which tells us if (and from which party: 1 for dems, -1 for reps) an incumbent is running. We want to find $\beta_0, \beta_1, \beta_2$.

Notice that we are regressing on a discrete variable I . This incumbency factor takes values 1, -1, or 0. As such it only changes the intercept of the regression. You can think of it as 3 regression lines, one for each subpopulation of incumbency, with their slope constrained to be the same. An intercept of β_0 for open seats, $\beta_2 + \beta_0$ for Democratic incumbents and $-\beta_2 + \beta_0$ for Republican incumbents.

You then think a little bit more and realize that, for example, in many conservative districts you will have a republican elected whether he/she is an incumbent or not. And you now realize that our analysis does not consider the party of the incumbent. So you decide to fix this

Lets define $P_{t,i}$ as the party in power right now before the election at time t , i.e. the party that won the election at time $t - 2$ in county i . It takes on values:

$$P = \begin{cases} -1 & \text{Republican Seat holder} \\ 1 & \text{Democratic Seat holder} \end{cases}$$

We can do this regression instead:

$DP1 \sim DP + I + P$, where

P represents the incumbent party, i.e. the party which won the election in year $t-2$.

In mathematical notation we have:

$$d_{t,i} = \beta_1 d_{t-2,i} + \beta_2 I_{t,i} + \beta_3 P_{t,i} + \beta_0,$$

where $P_{t,i}$ is the party in power right now before the election at time t , i.e. the party that won the election at time $t - 2$ in county i . The value of P is 1 for democrats, and -1 for republicans.

Interpretable Regressions

One can say that the coefficient of I now more properly captures the effect of incumbency, after controlling for party.

Regression coefficients become harder to interpret with multiple features. The meaning of any given coefficient depends on the other features in the model. Gelman and Hill advise: Typical advice is to interpret each coefficient "with all the other predictors held constant." [Gelman, Andrew; Hill, Jennifer (2006-12-25). Data Analysis Using Regression and Multilevel/Hierarchical Models] Economists like to use the phrase "ceteris paribus" to describe this.

The way to do this is interpretation to look at the various cases and explain what the co-efficients of P and I mean. Let us at first set I to 0 meaning no incumbents and explain what the coefficients of P mean. We are then fitting:

$$d_{t,i} = \beta_1 d_{t-2,i} + \beta_3 P_{t,i} + \beta_0,$$

which for the $P = 1$ (Democrat party winning the past election) case, gives us:

$$d_{t,i} = \beta_1 d_{t-2,i} + \beta_3 + \beta_0,$$

and, for the $P = -1$ (Republican party winning the past election) case, gives us:

$$d_{t,i} = \beta_1 d_{t-2,i} - \beta_3 + \beta_0.$$

You can see that β_3 then captures half the difference in the effect between democrats and republicans that comes from just having the party incumbent. It tells us that, with respect to the national swing measure β_0 , what's the party effect for republicans and democrats. It does it very poorly by splitting the difference between the democratic and republican party effects and being constant across seats, but its a start.

In [20]:

1	pairs=[
2	(1898,1896),
3	(1900,1898),
4	(1904,1902),
5	(1906,1904),
6	(1908, 1906),
7	(1910, 1908),
8	(1914, 1912),
9	(1916, 1914),
10	(1918, 1916),
11	(1920, 1918),
12	(1924, 1922),
13	(1926, 1924),
14	(1928, 1926),
15	(1930, 1928),
16	(1934, 1932),
17	(1936, 1934),
18	(1938, 1936),
19	(1940, 1938),
20	(1944, 1942),
21	(1946, 1944),
22	(1948, 1946),
23	(1950, 1948),
24	(1954, 1952),
25	(1956, 1954),
26	(1958, 1956),
27	(1960, 1958),
28	(1964, 1962),
29	(1966, 1964),
30	(1968, 1966),
31	(1970, 1968),
32	(1974, 1972),
33	(1976, 1974),
34	(1978, 1976),
35	(1980, 1978),
36	(1984, 1982),
37	(1986, 1984),
38	(1988, 1986),
39	(1990, 1988)
40]

Each CSV file has the following information:

- a number for the state
- a number for the district

- D1 and R1, the dem and repub percentages in the past election, and I1 the incumbency back then
- D and R, the dem and repub percentages in the present election, and I the incumbency now
- P, the incumbent party from the past election in that seat, 1 for democrats, -1 for republicans
- PNOW, the party which won the current election, 1 for democrats, -1 for republicans
- A variable we'll call T (for treatment), where we want to decide if we should replace an incumbent for a new candidate, or not.

$$T = \begin{cases} 0 & \text{Incumbent Running} \\ 1 & \text{New Candidate Running} \end{cases}$$

(This column is not used in this homework)

```
In [21]: 1 pairframes={}
2 for p in pairs:
3     key = str(p[0])+"-"+str(p[1])
4     pairframes[key] = pd.read_csv('data/housedata/{}.csv'.format(key))
```

To get warmed up, let us consider the 1988-1990 election pair.

```
In [22]: 1 pairframes['1990-1988'].head()
```

```
Out[22]:
```

	state	district	I	D	R	D1	R1	DP	DP1	P	PNOW	T
0	1	1	1	126566	50690	176463	51985	0.714029	0.772443	1	1	0.0
1	1	2	1	105085	70922	143326	81965	0.597050	0.636182	1	1	0.0
2	1	3	0	90772	83440	147394	74275	0.521043	0.664928	1	1	1.0
3	1	4	-1	32352	105682	55751	147843	0.234377	0.273834	-1	-1	0.0
4	1	5	0	85803	93912	58612	163729	0.477439	0.263613	-1	-1	1.0

To carry out the linear regression we'll use `statsmodels` from python, using the `ols`, or Ordinary Least Squares method defined there.

We use the `statsmodels` formula notation. `DP ~ DP1 + I` means linear regress `DP`, the democratic fraction of the vote this time around for a given house seat on `DP1` which is the democratic fraction the previous time around and `I`, a "factor" or categorical(nominal) variable with 3 levels:

```
In [23]: 1 import statsmodels.api as sm
2 from statsmodels.formula.api import glm, ols

/anaconda3/lib/python3.6/site-packages/statsmodels/compat/pandas.py:56: FutureWarning: The pandas.core.datetools module is deprecated and will be removed in a future version. Please use the pandas.tseries module instead.
from pandas.core import datetools
```

```
In [24]: 1 ols_model = ols('DP ~ DP1 + I', pairframes['1990-1988']).fit()
2 ols_model
```

```
Out[24]: <statsmodels.regression.linear_model.RegressionResultsWrapper at 0x11994d978>
```

```
In [25]: 1 ols_model.summary()
```

```
Out[25]: OLS Regression Results
```

Dep. Variable:	DP	R-squared:	0.806
Model:	OLS	Adj. R-squared:	0.804
Method:	Least Squares	F-statistic:	612.0
Date:	Wed, 26 Sep 2018	Prob (F-statistic):	1.04e-105
Time:	19:20:50	Log-Likelihood:	368.81
No. Observations:	298	AIC:	-731.6
Df Residuals:	295	BIC:	-720.5
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.2326	0.020	11.503	0.000	0.193	0.272
DP1	0.5622	0.040	14.220	0.000	0.484	0.640
I	0.0429	0.008	5.333	0.000	0.027	0.059

Omnibus:	7.465	Durbin-Watson:	1.728
Prob(Omnibus):	0.024	Jarque-Bera (JB):	7.316
Skew:	0.374	Prob(JB):	0.0258
Kurtosis:	3.174	Cond. No.	13.1

Interpretable Regressions

One can say that **The coefficient of `I` now more properly captures the effect of incumbency, after controlling for party.**

Regression coefficients become harder to interpret with multiple features. The meaning of any given coefficient depends on the other features in the model. Gelman and Hill advise: **Typical advice is to interpret each coefficient "with all the other predictors held constant."**[Gelman, Andrew; Hill, Jennifer (2006-12-25). Data Analysis Using Regression and Multilevel/Hierarchical Models] Economists like to use the phrase "ceteris paribus" to describe this.

The way to do this is interpretation to look at the various cases and explain what the co-efficients of P and I mean. Let us at first set I to 0 meaning no incumbents and explain what the coefficients of P mean. We are then fitting:

$$d_{t,i} = \beta_1 d_{t-2,i} + \beta_3 P_{t,i} + \beta_0,$$

which for the $P = 1$ (Democrat party winning the past election) case, gives us:

$$d_{t,i} = \beta_1 d_{t-2,i} + \beta_3 + \beta_0,$$

and, for the $P = -1$ (Republican party winning the past election) case, gives us:

$$d_{t,i} = \beta_1 d_{t-2,i} - \beta_3 + \beta_0.$$

You can see that β_3 then captures half the difference in the effect between democrats and republicans that comes from just having the party incumbent. It tells us that, with respect to the national swing measure β_0 , what's the party effect for republicans and democrats. It does it very poorly by splitting the difference between the democratic and republican party effects and being constant across seats, but its a start.

4.1 Explain the coefficient of Incumbency

Use a similar argument to the one above.

(Note that setting I to 1 also constrains P to 1, but the reverse is not true as we saw above).

4.1 Answer

When we are fitting: $d_{t,i} = \beta_1 d_{t-2,i} + \beta_2 I_{t,i} + \beta_3 P_{t,i} + \beta_0$,

- $I = -1$, the incumbency also constraints $P = -1$:

$$d_{t,i} = \beta_1 d_{t-2,i} - \beta_2 - \beta_3 + \beta_0$$

- $I = 1$, the incumbency also constraints $P = 1$:

$$d_{t,i} = \beta_1 d_{t-2,i} + \beta_2 + \beta_3 + \beta_0$$

- $I = 0$, when incumbency does not have constraint on P :

- $P = -1$:

$$d_{t,i} = \beta_1 d_{t-2,i} - \beta_3 + \beta_0$$

- $P = 1$:

$$d_{t,i} = \beta_1 d_{t-2,i} + \beta_3 + \beta_0$$

Therefore, the coefficient of I , β_2 , captures the effect of incumbency, after controlling for the effect of party which is captured by the coefficient of P , β_3 .

4.2 Carry out the linear regression $DP_t = \alpha_0 + \alpha_1 I_t + \alpha_2 P_t$ for all the year pairs

Present the results in a dataframe `ols_frame`. Comment on the trend in the incumbency coefficients after 1960.

(FORMAT: This dataframe has columns `yp`, the year-pair string (the keys of the dictionary of frames), the year for which we do the regression `year` (the higher year in the pair), the `formula`, which is just repeated, and the R-squared in `R2` for each regression, as well as the parameters of the regression and the p-values for the regression (for the name of the column here prefix the parameter with `p_` to denote the p-value).)

4.2 Answer - Codes

```

In [26]: 1 fm = 'DP ~ DP1 + I + P'
2 ols_result_dict_list = []
3 for yp_key, yp_df in pairframes.items():
4     result_dict = {}
5
6     # fit OLS
7     yp_ols = ols(fm, yp_df).fit()
8
9     # store OLS results
10    result_dict['yp'] = yp_key
11    result_dict['year'] = int(yp_key.split('-')[0])
12    result_dict['formula'] = fm
13    result_dict['R2'] = yp_ols.rsquared
14    result_dict.update(yp_ols.params.to_dict()) # fitted intercept and coefficients
15    result_dict.update(yp_ols.pvalues.add_prefix('p_').to_dict()) # pvalues
16
17    # append the old result dictionary to the list
18    ols_result_dict_list.append(result_dict)
19
20 # convert the list of dictionaries to `ols_frame`
21 ols_frame = pd.DataFrame(ols_result_dict_list)
22 ols_frame = ols_frame[['yp', 'year', 'formula', 'R2', 'Intercept', 'DP1', 'I', 'P', 'p_Intercept', 'p_DP1', 'p_I', 'p_P']]
23 print(ols_frame.shape)
24 ols_frame.head()

```

(38, 12)

```

Out[26]:

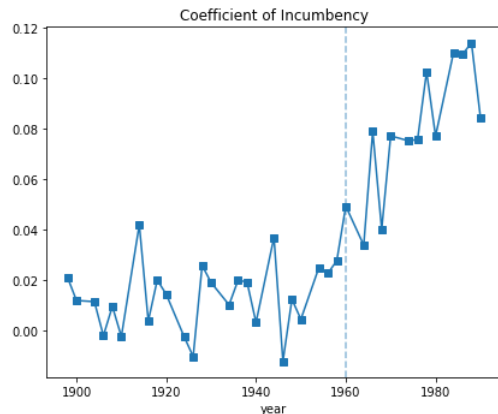
```

	yp	year	formula	R2	Intercept	DP1	I	P	p_Intercept	p_DP1	p_I	p_P
0	1898-1896	1898	DP ~ DP1 + I + P	0.714405	0.091247	0.901581	0.021063	-0.006020	3.412383e-03	1.286508e-33	0.035734	0.558655
1	1900-1898	1900	DP ~ DP1 + I + P	0.819429	0.098974	0.768643	0.011889	0.002537	3.999435e-07	9.145154e-58	0.154437	0.747252
2	1904-1902	1904	DP ~ DP1 + I + P	0.867082	-0.005676	0.924338	0.011397	0.001106	7.656457e-01	7.903248e-78	0.082649	0.871755
3	1906-1904	1906	DP ~ DP1 + I + P	0.856573	0.098251	0.882225	-0.002075	0.017502	3.586958e-07	1.928989e-72	0.782880	0.029587
4	1908-1906	1908	DP ~ DP1 + I + P	0.863811	0.103591	0.778613	0.009547	-0.003617	2.381739e-11	3.310047e-85	0.229116	0.665677

```

In [27]: 1 # inspect the fitted coefficient of Incumbency after year 1960
2 fig, ax = plt.subplots(1, 1, figsize=(6, 5))
3 ols_frame.set_index('year')['I'].plot(style='s-', title='Coefficient of Incumbency', ax=ax)
4 ax.axvline(1960, linestyle='--', alpha=0.5)
5 plt.tight_layout()

```



4.2 Answer The coefficients of incumbency started to increase after year 1960.

4.3 Bootstrap a distribution for the coefficient of I for 1990-1988

Plot a histogram of the distribution of the co-efficient. Also print the the 2.5th and 97.5th quantile of the distribution to give a non-parametric confidence interval, plotting these on the histogram. What conclusions can you draw?

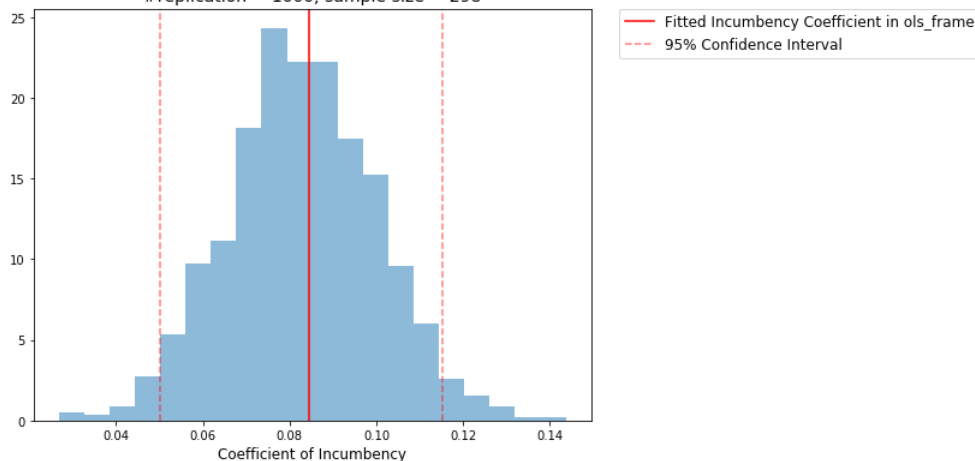
(Hint: Bootstrap involves sampling with replacement from the data and recalculating the quantity of interest, in our case the regression. This will give you a new coefficient for each regression. If you're interested in using the method for more complex applications it if imperative to familiarize with the assumptions, [this](#) (<http://stats.stackexchange.com/questions/26088/explaining-to-laypeople-why-bootstrapping-works>) is a good start, but [this article \(https://goo.gl/2T6k8j\)](https://goo.gl/2T6k8j) is also helpful.)

```
In [28]: 1 np.random.seed(99)
2
3 # bootstrap data samples
4 fm = 'DP ~ DP1 + I + P'
5 nSim = 1000
6 sampleSize = pairframes['1990-1988'].shape[0]
7
8 I_coef_list = []
9 for i in range(nSim):
10     sample_idx = np.random.choice(pairframes['1990-1988'].index.values, size=(sampleSize,), replace=True)
11     sample_df = pairframes['1990-1988'].iloc[sample_idx]
12
13     # fit OLS
14     sample_df_ols = ols(fm, sample_df).fit()
15     I_coef_list.append(sample_df_ols.params['I'])
16
17 I_coef_pct = np.percentile(I_coef_list, [2.5, 97.5])
18 print('The 95% confidence interval: [{}, {}]' .format(I_coef_pct[0], I_coef_pct[1]))
```

The 95% confidence interval: [0.05015926585810058, 0.1153782022427402]

```
In [29]: 1 # plot the I's coefficients on bootstrapped samples
2 fig, ax = plt.subplots(1, 1, figsize=(8, 6))
3 ax.hist(I_coef_list, density=True, bins=20, alpha=0.5)
4 ax.axvline(ols_frame.set_index('year').loc[1990]['I'], c='r', label='Fitted Incumbency Coefficient in ols_frame')
5 ax.axvline(I_coef_pct[0], c='r', linestyle='--', alpha=0.5)
6 ax.axvline(I_coef_pct[1], c='r', linestyle='--', alpha=0.5, label='95% Confidence Interval')
7 ax.set_title('Histogram (density) of Bootstrapped Coefficients of Incumbency 1990-1988 \n#replication = {}, sample size = {}'.format(nSim, sampleSize))
8 ax.set_xlabel('Coefficient of Incumbency', fontsize=12)
9 plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0., fontsize=12)
10 plt.tight_layout()
```

Histogram (density) of Bootstrapped Coefficients of Incumbency 1990-1988
#replication = 1000, sample size = 298



4.3 Answer - Conclusions

The coefficients of Incumbency on the 1000x bootstrapped samples have a confidence interval of [0.05, 0.12], not containing zero, therefore suggesting Incumbency being a significant predictor.

4.4 Inference using p-values over time

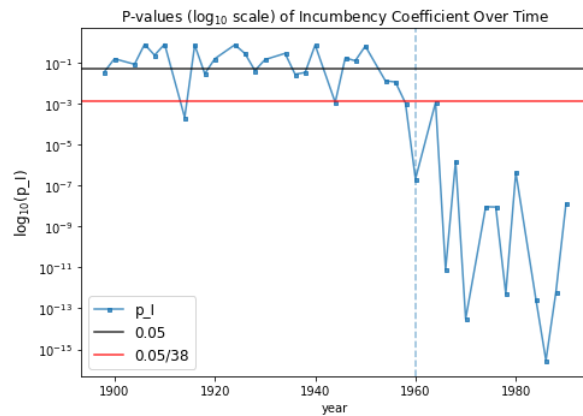
Of-course, another more classic way of doing this same inference is through the regression itself -- it gives us p-values. These are values from a t-test that asks if the coefficient is different from 0. The regression machinery assumes Normality of errors for this purpose. Let's assume the Normality and do an inference on all the years in our regression. The assumption used to calculate these p-values are: for each model (in our case year), the errors at each point of the regression are uncorrelated and follow a Normal distribution. We shall assume these to be true for now (in real life you ought to be checking a plot of residuals as well).

Generally we'd like the p-values to be vanishingly small as they represent the probability that we observed such an extreme incumbency effect purely by chance. Have a look at the Wikipedia page on [p-values \(https://en.wikipedia.org/wiki/P-value\)](https://en.wikipedia.org/wiki/P-value) for a quick reminder.

Furthermore, when constructing results like this (where there are many tests considered at once) there are other concerns to take into account. One such concern is the issue of [multiple testing \(https://en.wikipedia.org/wiki/Multiple_comparisons_problem\)](https://en.wikipedia.org/wiki/Multiple_comparisons_problem). This is important because when we start dealing with a **large number** of hypotheses jointly the probability of making mistakes gets larger, hence we should be **more stringent** about what it means for a result to be significant. One such correction is the [Bonferroni Correction \(https://en.wikipedia.org/wiki/Bonferroni_correction\)](https://en.wikipedia.org/wiki/Bonferroni_correction) which provides a new bound for deciding significance. Instead of asking the classic question: **is the p-value < 0.05?**, this considers instead a stricter bound, we ask: **is p-value < 0.05/H**. Where H is the number of hypotheses being considered, in our case $H = 38$ (the number of years) -- this is a much higher bar for significance.

Plot a graph of incumbency (α) coefficient p-values for every year. Use this plot to study if the coefficients after 1960 are significantly different from 0. (Plot them in log scale for easier viewing of small numbers. Also draw lines at $\log(0.05)$ and $\log(0.05/38)$ for reference). Interpret your results.

```
In [30]: 1 fig, ax = plt.subplots(1, 1, figsize=(7, 5))
2 ols_frame.set_index('year')['p_I'].plot(style='s-', logy=True, ax=ax, ms=3, alpha=0.8,
3                                           title=r'P-values ( $\log_{10}$  scale) of Incumbency Coefficient Over Time')
4 # ols_frame.plot(kind='scatter', x='year', y='p_I', logy=True, marker='s', ax=ax)
5 ax.axvline(1960, linestyle='--', alpha=0.5)
6 ax.axhline(0.05, c='k', linestyle='-', alpha=0.8, label='0.05')
7 ax.axhline(0.05/38, c='r', linestyle='-', alpha=0.8, label='0.05/38')
8 ax.set_ylabel(r' $\log_{10}(p_I)$ ', fontsize=12)
9 plt.legend(fontsize=12)
10 plt.tight_layout()
11
```



4.4 Answer - Conclusions

It is observed that the p-values of Incumbency coefficients were below 0.05/38 after year 1960. This suggests that under comparisons with a total 38 hypotheses, Incumbency has become a significant predictor after 1960.

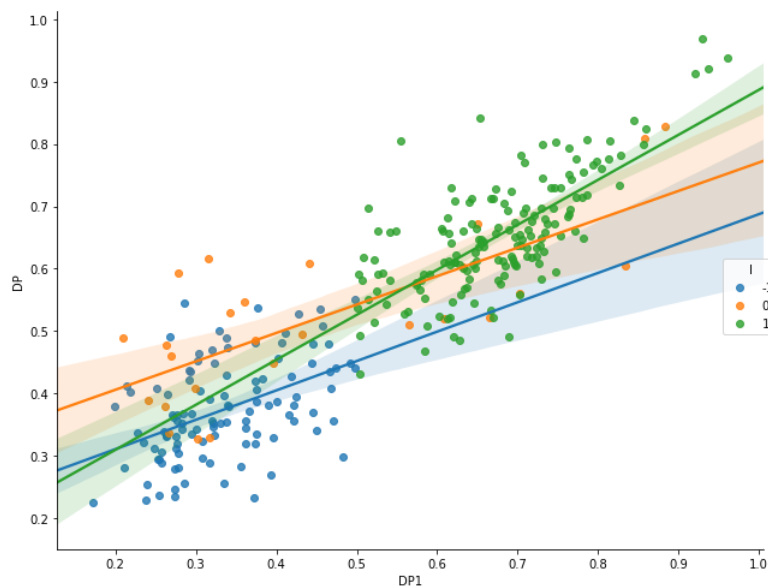
4.5 Carry out the linear regression with an interaction between the previous elections democratic fraction and this elections incumbency, for all the year pairs

Is the regression complete? Or do we need more features?

Recall that our model is fairly restrictive, the different incumbency groups are allowed to have different intercepts but the new candidate group, $I = 0$ is equally between the two incumbency groups. Furthermore, the incumbency groups are not allowed different slopes, meaning the effect of the previous elections Democratic fraction ($DP1$) is assumed the same for all incumbency groups. This may not be the case.

In the figure below we can see that in fact the different groups seem to have not only different intercepts, but also possibly different slopes.

```
In [31]: 1 sns.lmplot(x="DP1", y="DP", hue = "I", data=pairframes['1990-1988'], size = 7, aspect=1.2)
2 plt.tight_layout()
```



Carry out the regression with an between the previous elections democratic fraction and this elections incumbency, for each year pair. Is there evidence for interaction? How can you know for sure?

(HINT: In mathematical notation this regression is:

$$d_{i,t} = \beta_4 d_{i,t-2} I_{i,t} + \beta_3 I_{i,t} + \beta_2 P_{i,t} + \beta_4 I_{i,t} d_{i,t-2} + \beta_0,$$

In statsmodels notation, we wish to carry out the regression:

DP ~ DP1 + I + P + DP1:I)

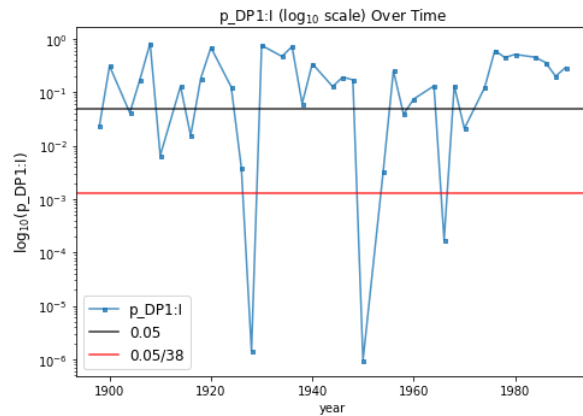
```
In [32]: 1 # regression with interaction terms
2 fm_int = 'DP ~ DP1 + I + P + DP1:I'
3 ols_int_result_dict_list = []
4 for yp_key, yp_df in pairframes.items():
5     int_result_dict = {}
6
7     # fit OLS
8     yp_ols_int = ols(fm_int, yp_df).fit()
9
10    # store OLS results
11    int_result_dict['yp'] = yp_key
12    int_result_dict['year'] = int(yp_key.split('-')[0])
13    int_result_dict['formula'] = fm_int
14    int_result_dict['R2'] = yp_ols_int.rsquared
15    int_result_dict.update(yp_ols_int.params.to_dict()) # fitted intercept and coefficients
16    int_result_dict.update(yp_ols_int.pvalues.add_prefix('p_').to_dict()) # pvalues
17
18    # append the old result dictionary to the list
19    ols_int_result_dict_list.append(int_result_dict)
20
21 # convert the list of dictionaries to `ols_int_frame`
22 ols_int_frame = pd.DataFrame(ols_int_result_dict_list)
23 ols_int_frame = ols_int_frame[
24     ['yp', 'year', 'formula', 'R2', 'Intercept', 'DP1', 'I', 'P', 'DP1:I', 'p_Intercept', 'p_DP1', 'p_I', 'p_P', 'p_DP1:I']
25 ]
26 print(ols_int_frame.shape)
27 ols_int_frame.head()
28
```

(38, 14)

```
Out[32]:
```

	yp	year	formula	R2	Intercept	DP1	I	P	DP1:I	p_Intercept	p_DP1	p_I	p_P	p_DP1:I
0	1898-1896	1898	DP ~ DP1 + I + P + DP1:I	0.721095	0.062006	0.939910	-0.041304	-0.014577	0.142585	6.246192e-02	2.636899e-34	0.156726	0.180435	0.023492
1	1900-1898	1900	DP ~ DP1 + I + P + DP1:I	0.820178	0.097581	0.765434	-0.006918	0.002967	0.036985	6.079301e-07	3.920054e-57	0.733069	0.706569	0.309648
2	1904-1902	1904	DP ~ DP1 + I + P + DP1:I	0.868971	0.002251	0.896128	-0.025120	0.002800	0.073879	9.072736e-01	1.655357e-69	0.186036	0.683325	0.041088
3	1906-1904	1906	DP ~ DP1 + I + P + DP1:I	0.857476	0.093900	0.881553	-0.023508	0.016063	0.046471	1.479223e-06	2.130199e-72	0.178058	0.047275	0.173509
4	1908-1906	1908	DP ~ DP1 + I + P + DP1:I	0.863833	0.103742	0.779571	0.012644	-0.003498	-0.006522	2.515951e-11	6.057820e-84	0.433447	0.677127	0.825486

```
In [33]: 1 # plot p-values of the interaction term
2 fig, ax = plt.subplots(1, 1, figsize=(7, 5))
3 ols_int_frame.set_index('year')['p_DP1:I'].plot(style='s-', logy=True, ax=ax, ms=3, alpha=0.8,
4         title=r'p_DP1:I ($\log_{10}$ scale) Over Time')
5
6 # ax.axvline(1960, linestyle='--', alpha=0.5)
7 ax.axhline(0.05, c='k', linestyle='-', alpha=0.8, label='0.05')
8 ax.axhline(0.05/38, c='r', linestyle='-', alpha=0.8, label='0.05/38')
9 ax.set_ylabel(r'$\log_{10}$ (p_DP1:I)', fontsize=12)
10 plt.legend(fontsize=12)
11 plt.tight_layout()
12
```



```
In [34]: 1 # print years where p_DP1:I < 0.05/38
2 ols_int_frame[ols_int_frame['p_DP1:I'] < 0.05/38][['yp', 'year', 'p_DP1:I']]
```

```
Out[34]:
```

	yp	year	p_DP1:I
12	1928-1926	1928	1.425443e-06
21	1950-1948	1950	9.623747e-07
27	1966-1964	1966	1.715822e-04


```

In [35]: 1 # bootstrap a year where DP1:I is SIGNIFICANT
2 np.random.seed(99)
3
4 yp_sig = ols_int_frame[ols_int_frame['p_DP1:I'] < 0.05/38]['yp'].iloc[0]
5 nSim = 1000
6 sampleSize = pairframes[yp_sig].shape[0]
7
8 DP1xI_coef_list = []
9 for i in range(nSim):
10     sample_idx = np.random.choice(pairframes[yp_sig].index.values, size=(sampleSize,), replace=True)
11     sample_df = pairframes[yp_sig].iloc[sample_idx]
12
13     # fit OLS
14     sample_df_ols = ols(fm_int, sample_df).fit()
15     DP1xI_coef_list.append(sample_df_ols.params['DP1:I'])
16
17 DP1xI_coef_pct = np.percentile(DP1xI_coef_list, [2.5, 97.5])
18 print('The 95% confidence interval during {}: {}'.format(yp_sig, DP1xI_coef_pct[0], DP1xI_coef_pct[1]))

```

The 95% confidence interval during 1928-1926: [0.0945393307337304, 0.2584001729589221]

```

In [36]: 1 # bootstrap a year where DP1:I is NOT SIGNIFICANT
2 np.random.seed(99)
3
4 yp_no_sig = ols_int_frame[ols_int_frame['p_DP1:I'] > 0.1]['yp'].iloc[0]
5 nSim = 1000
6 sampleSize = pairframes[yp_no_sig].shape[0]
7
8 DP1xI_coef_list = []
9 for i in range(nSim):
10     sample_idx = np.random.choice(pairframes[yp_no_sig].index.values, size=(sampleSize,), replace=True)
11     sample_df = pairframes[yp_no_sig].iloc[sample_idx]
12
13     # fit OLS
14     sample_df_ols = ols(fm_int, sample_df).fit()
15     DP1xI_coef_list.append(sample_df_ols.params['DP1:I'])
16
17 DP1xI_coef_pct = np.percentile(DP1xI_coef_list, [2.5, 97.5])
18 print('The 95% confidence interval during {}: {}'.format(yp_no_sig, DP1xI_coef_pct[0], DP1xI_coef_pct[1]))

```

The 95% confidence interval during 1900-1898: [-0.09161739278835825, 0.15939972848665207]

4.5 Answer - Evidence of the Interaction Term

- Evidence from p-values

As shown above, only 3 years (1928, 1950, 1966) have $p_{DP1:I} < 0.05/38$. Therefore, only these 3 years but not the others have significant evidence for the interaction term $DP1:I$.

- Evidence from the confidence interval of coefficient estimates on the bootstrapped samples not containing zero

Consistent with p-values $< 0.05/38$, the 95% Confidence Interval (CI) of $DP1:I$'s coefficient estimates on a 1000x non-parametric bootstrapping does not contain zero for the years: but does contain zeros for the others.

As an example year with significant evidence of $DP1:I$ ($p_{DP1:I} < 0.05/38$), 1928-1926, the CI is [0.09, 0.26], not containing 0 and thus suggesting that $DP1:I$ is a significant predictor during this year pair.

Similarly, as an example year without significant evidence of $DP1:I$ ($p_{DP1:I} > 0.1$), 1900-1898, the CI is [-0.09, 0.16], containing 0 and thus suggesting that $DP1:I$ is not significant during this year pair.

In []: 1